CS 547: Assignment 04

Programming Assignments (75%)

The goal of this assignment is to write code to **train and evaluate PyTorch neural network models to predict classes from the CIFAR10 dataset.**

You must train and evaluate at least FOUR different variations.

A04.py

You will write the following functions in this python script:

- **get_approach_names()**
- **get_approach_description(approach_name)**
- **get_data_transform(approach_name, training)**
- **get_batch_size(approach_name)**
- **create_model(approach_name, class_cnt)**
- **train_model(approach_name, model, device, train_dataloader, test_dataloader)**

Note that most of these functions rely on the approach_name parameter. This is how you will code different behavior per approach.

get_approach_names()

Returns a list of the names of all combinations you will be testing. It is largely up to you what names you use, but self-documenting names are encouraged.

A rather boring example: ["CNN0", "CNN1"]

get_approach_description(approach_name)

Given the approach_name, return a text description of what makes this approach distinct. This does not have to be very long: one sentence is sufficient.

get_data_transform(approach_name, training)

Given the approach_name and whether this is for training data, return the appropriate dataset transform. You may add data augmentation if you like for training data transforms, BUT:

- Do NOT do data augmentation for non-training data!
- At minimum, you must perform the following (depending on your version of torchvision):

(0.15) data_transform = v2.Compose([v2.ToImageTensor(), v2.ConvertImageDtype()])
(0.16) data_transform = v2.Compose([v2.ToImage(), v2.ToDtype(torch.float32, scale=True)])

get_batch_size(approach_name)

Given the approach_name, return the preferred batch size. This can be hardcoded to a single value if you do not have a preference, BUT this allows you to use a smaller batch size if your architecture is large enough where that is a concern.

create_model(approach_name, class_cnt)

Given the approach_name and output class_cnt, build and return a PyTorch neural network that takes a color 32 x 32 image and outputs a vector with class_cnt elements.

You do not have to move this model to the GPU (that is done for you in the provided scripts).

Any of your networks should have the following attributes:

1. The input channels should be 3.
2. The output layer should have class_cnt nodes.
3. The network must have at least 3 convolutional and/or fully-connected layers, not including the last output layer.
4. **You are free to use ANY of the layer types available in PyTorch**.
5. You may choose to use one of the existing pre-trained networks (like VGGNet) and fine-tune it.

train_model(approach_name, model, device, train_dataloader, test_dataloader)

Given the provided model, the device it is located, and the relevant dataloaders, train this model and return it.

A couple of caveats:

- **You can ONLY train on the data from train_dataloader.**
- You may print out how the network is doing on test_dataloader. **However, you may NOT use test_dataloader as a validation set** (e.g., you cannot use it to tune learning rates dynamically).
- You are free to choose any optimizer you wish.
- You are free to decide how many epochs of training you use.
- **Your loss SHOULD be CrossEntropyLoss.** You are, however, welcome to add other losses if you believe it will help.

Note that, in general, differences here do NOT count as acceptable approach differences, such as:

- **Epoch counts**
- **Different optimizers**

Acceptable Approach Differences

Your approaches should be sufficiently different from each other! Acceptable differences include:

- Different sizes of filters
- Notably different filter counts (32 vs. 64, not 32 vs. 33)
- Notably different node counts
- Consistently different activation functions.
- Adding 1 or more additional layers
- Adding skip connections
- With or without significant data augmentation (or sufficiently different variations of data augmentation)

I actively encourage you to look for examples of working networks.

HOWEVER, if you do use code from another source:

- **Cite it in the body of the code like so:**

 # Taken from: https://keras.io/getting-started/sequential-model-guide/

- **Include it in your README.md**

NOTE: Even if you cite it, you may NOT use code from another student (whether presently in the course or from a previous semester).

However, any code in the slides for this course is freely usable.

If you are unsure whether any variants you intend to experiment with are different enough, ASK ME.

TRAINING CAN BE VERY TIME-CONSUMING. DO NOT LEAVE THIS UNTIL LAST MINUTE!

Results Folder (15%)

Make sure the contents of **assign04/output** are committed, including:

- **Your model files (model_*.pth) UNLESS they are over 15MB!**
- **ALL_RESULTS.txt**

PLEASE add the following to your .gitignore file:

- data/
    - This is where CIFAR10 is downloaded
- assign04/output/bigmodel.pth
    - Any models over 15MB

README (10%)

Update the README.md of your repository with the following:

- Add a new entry under "Runnable Python Scripts" for "**A04.py**". This should contain the following subsections:
    - ***Methods:*** Briefly describe your different methods.
    - ***Results:*** Include your training and testing accuracy and F1 scores in a table.
    - ***Trained Models:***
        - If your saved model(s) are larger than 15MB, include instructions on how to download the models. You may use external storage sites to provide the models (e.g., OneDrive).
        - If everything is already in the repo, indicate that.
    - ***References:*** Any citations to external network architectures used (any reasonable citation method is allowed).

Provided Programs

Train_A04.py

- Asks you which approach you wish to train.
- Loads the data.
- Trains the model.
- Saves the model to assign04/output

Eval_A04.py

- Asks you which approach you wish to evaluate (if -1, evaluates all).
- Evaluates the model(s) on training and testing data.
- Prints and saves the results for each approach to assign04/output with filename: <approach name>_RESULTS.txt or ALL_RESULTS.txt.

BONUS POINTS

For the entire assignment, bonus points will be awarded to the submission that has the following:

- (+%5) Best **testing F1 score** in the class
- (+%5) Experimenting with 8 variants or more

Grading

- (75%) Programming
- (15%) Output/results folder
- (10%) README
- (+10%) Bonus credit

I reserve the right to take points off for not meeting the specifications in this assignment description.

The following specific penalties (applied to the WHOLE assignment) can be expected:

Problem	**Maximum Penalty**
Code that is not syntactically correct	60
Less than 4 approaches	40
Approaches not sufficiently distinct	30
Pretrained models not available (either directly or through download)	30
Missing output/results folder	15
README update missing/incorrect	10
Technically correct code, but poor code design	5
Missing assignment Git branch	5

In general, these are things that will be penalized:

- **Code that is not syntactically correct (up to 60 points off!)**
- Sloppy or poor coding style
- Bad coding design principles
- Code that crashes, does not run, or takes a VERY long time to complete
- Using code from ANY source other than the course materials
- Collaboration on code of ANY kind; this is an INDIVIDUAL PROJECT
- Sharing code with other people in this class or using code from this or any other related class
- Output that is incorrect
- Algorithms/implementations that are incorrect
- Submitting improper files
- Failing to submit ALL required files